

Closing the Triage Loop

MCP-mediated KLEE validation of
LLM-reported C/C++ vulnerabilities

Christopher Scherb, Alexander Trapp, Ruben Hutter,
Nico Bachmann, Luc Bryan Heitz

University of Applied Sciences and Arts Northwestern Switzerland

KES 2026 | **Dublin**

curl: AI reports created real work for maintainers

< 5%

of security submissions
confirmed as vulnerabilities
in 2025

- ▶ Low-quality AI reports consumed maintainer time
- ▶ curl ended bounty rewards on **31 January 2026**
- ▶ By April, AI-assisted reports had improved in quality; volume was higher

A convincing report still needs independent verification

Daniel Stenberg: The end of the curl bug-bounty (26 Jan 2026); High-Quality Chaos (22 Apr 2026).

The problem: which AI findings are real?

**An LLM flags a possible C/C++ bug.
Would you report it?**

- ▶ A convincing explanation is not execution evidence
- ▶ A failure may depend on inputs that real callers never provide
- ▶ Manually building a test and checking callers takes work

Our goal: check each candidate and explain the conditions under which it can fail

The LLM selects what to check


- ▶ Read the code and identify a suspicious function
- ▶ Identify inputs that should vary during analysis
- ▶ Inspect callers and propose assumptions about valid inputs
- ▶ Submit a structured candidate to the verification tool



The LLM supplies a hypothesis and context; execution must check it


Symbolic execution checks possible executions

- ▶ Use variables instead of one fixed test input
- ▶ Follow branches and collect their conditions
- ▶ Solve those conditions to find possible paths
- ▶ If a path fails, obtain a concrete input that reaches it




KLEE is the symbolic execution engine used in our system

The bridge: a small program KLEE can execute



- ▶ A **harness** sets up inputs and calls the selected function
- ▶ The server turns declared inputs into symbolic variables
- ▶ Assumptions restrict those variables before the call
- ▶ The harness and target code are compiled to LLVM bitcode

How the LLM and KLEE work together




- ▶ **LLM**: choose the candidate, inputs, and assumptions
- ▶ **MCP server**: generate the harness and manage KLEE runs
- ▶ **KLEE**: return a verdict and execution evidence
- ▶ **LLM**: explain the result and investigate caller context

What the verdicts mean

Confirmed	A matching error was found; a concrete input demonstrates it
Infeasible	No matching error exists in the fully explored input model
Unresolved	The run timed out, failed to build, or hit an engine error

- ▶ **Confirmed:** check whether real callers can supply that input
- ▶ **Infeasible:** ask which inputs the assumptions excluded


A caller contract is a promise about inputs



- ▶ **Caller:** the code that invokes the function
- ▶ **Contract:** “I will pass at most four records”
- ▶ The LLM represents that promise as the bound $0 \leq n \leq 4$

An assumption in the harness does not add a check to the real function


Bounds remove paths from the search



- ▶ Paths requiring $n > 4$ conflict with the bound and are discarded
- ▶ In a loop over n records, at most four iterations are needed
- ▶ Fewer execution states usually mean less solver work and memory

Less exploration also means a narrower safety claim

Why retry: what if the function is used differently?



Possible results: an error appears; no error is found after completion; or the larger search remains unresolved.

- ▶ The LLM may have inferred a bound that is too restrictive
- ▶ A new caller or refactor may pass more records, such as $n = 8$
- ▶ Retry without the LLM bounds to check excluded behavior

A new error can reveal a fragile assumption, even if today's usage is safe


Proof-of-concept evaluation

14

hand-constructed
candidates

- ▶ Toy programs for selected weakness classes
 - ▶ Bounds and retry sensitivity cases
 - ▶ Extracted libpng functions
 - ▶ Whole-library libpng + zlib analysis
-
- ▶ KLEE 3.1; LLVM/Clang 13; STP + MiniSat
 - ▶ Whole-library target: 2.2 MB bitcode, 30 translation units

All 14 rows match their expected verdict




- ▶ 9 confirmed
- ▶ 5 infeasible
- ▶ 0.9–42.9 s per row

Evidence of internal consistency on these selected cases

- ▶ External precision, recall, and analyst time savings remain unmeasured


Bounds can make exploration finish



- ▶ Same libpng parser and harness; with and without a semantic bound
- ▶ The terminating result applies only under the added contract


The same function can be safe under its contract

- ▶ CVE-2015-8540 case study: an internal libpng helper
- ▶ Function-level analysis reports an error across three versions
- ▶ Adding the caller contract changes the verdict to infeasible



Fixing the caller can make normal usage safe while the helper still trusts it

Can today's callers reach the error?




A local error needs caller context

Reaching a function and reaching its error state are different questions.

- ▶ `png_format_number`: an unchecked internal precondition
- ▶ Shipped callers satisfy the precondition at every reviewed call site
- ▶ Hardening opportunity: guard the precondition against future misuse

Conclusion: the LLM guides, KLEE checks



- ▶ Focus analysis on a candidate and the way it is called
- ▶ Use bounds to make that analysis tractable
- ▶ Retry to reveal behavior hidden by the assumptions

Report the verdict together with the assumptions that made it possible

Would you report it now?

- ▶ **Evidence:** can an execution demonstrate the error?
- ▶ **Assumptions:** for which inputs was the function checked?
- ▶ **Usage:** can a current or future caller violate those assumptions?

Check the finding. Explain its scope.

Questions?

github.com/FHNW-Security-Lab/klee-mcp

Backup: the benchmark rows

Candidate	Verdict	Wall time (s)
bof_01	confirmed	1.0
intoverflow_01	confirmed	2.4
null_01	confirmed	0.9
uaf_01	confirmed	0.9
safe_01	infeasible	1.1
bof_01_bounded	confirmed	0.9
bof_01_tootight	confirmed after retry	0.9
libpng_fp_number	infeasible	30.9
libpng_fp_number_bug	confirmed	5.9
libpng_fp_number_unb.	confirmed	9.4
libpng_sig_cmp	infeasible	5.1
libpng_fp_whole	infeasible	42.9
libpng_push_read	infeasible	4.1
libpng_format_number	confirmed; contract violation	4.2

Infeasible verdicts are conditional on the harness and engine model.

Backup: larger targets expose model limits

Target	Run result	Interpretation
libpng tRNS	240 s; timeout	Analysis unresolved
libjpeg-turbo	5.4 s; infeasible 1 completed path	Error handling limits exploration; weak evidence of safety
libxml2	23 s; infeasible 11 completed paths	A second whole-library run; still only the modeled inputs

- ▶ Whole-library linking is not whole-program coverage
- ▶ Read path statistics and unsupported behavior with the verdict

Backup: limitations and next steps

- ▶ **Small benchmark:** evaluate an independent set of real findings
- ▶ **Engine limits:** improve handling of complex parser error paths
- ▶ **Caller assumptions:** validate them against realistic entry points
- ▶ **Impact assessment:** review automated labels before reporting

A local failure is one step toward a useful security report

Backup: further reading

- ▶ Cadar, Dunbar, Engler (2008)
KLEE: Unassisted and Automatic Generation of High-Coverage Tests for Complex Systems Programs, OSDI
- ▶ Ramos and Engler (2015)
Under-Constrained Symbolic Execution: Correctness Checking for Real Code, USENIX Security

Project: github.com/FHNW-Security-Lab/klee-mcp